

# ARM Architecture Overview

## Development of the ARM Architecture

- Processor Architecture = Instruction Set + Programmer's model



**ARM7TDMI**  
**ARM922T**  
Thumb  
instruction set



**ARM926EJ-S**  
**ARM946E-S**  
**ARM966E-S**  
Improved  
ARM/Thumb  
Interworking  
DSP instructions  
**Extensions:**  
Jazelle (5TEJ)



**ARM1136JF-S**  
**ARM1176JZF-S**  
**ARM11 MPCore**  
SIMD Instructions  
Unaligned data support  
**Extensions:**  
Thumb-2 (6T2)  
TrustZone (6Z)  
Multicore (6K)



**Cortex-A8/R4/M3/M1**  
Thumb-2  
**Extensions:**  
v7A (applications) – NEON  
v7R (real time) – HW Divide  
V7M (microcontroller) – HW  
Divide and Thumb-2 only

- Note: Implementations of the same architecture can be very different
  - ARM7TDMI - architecture v4T. Von Neuman core with 3 stage pipeline
  - ARM920T - architecture v4T. Harvard core with 5 stage pipeline and MMU

# ARM Architecture profiles

- Application profile (ARMv7-A → e.g. Cortex-A8)
  - Memory management support (MMU)
  - Highest performance at low power
    - Influenced by multi-tasking OS system requirements
  - TrustZone and Jazelle-RCT for a safe, extensible system
- Real-time profile (ARMv7-R → e.g. Cortex-R4)
  - Protected memory (MPU)
  - Low latency and predictability 'real-time' needs
  - Evolutionary path for traditional embedded business
- Microcontroller profile (ARMv7-M → e.g. Cortex-M3)
  - Lowest gate count entry point
  - Deterministic and predictable behavior a key priority
  - Deeply embedded use

# Programmer's Model



# Data Sizes and Instruction Sets


- When used in relation to the ARM:
  - Halfword** means 16 bits (two bytes)
  - Word** means 32 bits (four bytes)
  - Doubleword** means 64 bits (eight bytes)
- Most ARMs implement two instruction sets
  - 32-bit **ARM** Instruction Set
  - 16-bit **Thumb** Instruction Set
- Latest ARM cores introduce a new instruction set **Thumb-2**
  - Provides a mixture of 32-bit and 16-bit instructions
  - Maintains code density with increased flexibility
- Jazelle-DBX cores can also execute **Java bytecode**

# Processor Modes

- The ARM has seven basic operating modes:
  - Each mode has access to **own** stack and a **different** subset of registers
  - Some operations can only be carried out in a privileged mode

Exception modes	Mode	Description	Privileged modes
	Supervisor (SVC)	Entered on reset and when a Software Interrupt instruction (SWI) is executed	
	FIQ	Entered when a high priority (fast) interrupt is raised	
	IRQ	Entered when a low priority (normal) interrupt is raised	
	Abort	Used to handle memory access violations	
	Undef	Used to handle undefined instructions	
	System	Privileged mode using the same registers as User mode	Unprivileged mode
	User	Mode under which most Applications / OS tasks run	

# The ARM Register Set




THE ARCHITECTURE FOR THE DIGITAL WORLD®

7

ARM®

# Program Status Registers



- **Condition code flags**
  - N = **N**egative result from ALU
  - Z = **Z**ero result from ALU
  - C = ALU operation **C**arried out
  - V = ALU operation **o**Verflowed
- **Sticky Overflow flag - Q flag**
  - Architecture 5TE and later only
  - Indicates if saturation has occurred
- **J bit**
  - Architecture 5TEJ and later only
  - J = 1: Processor in Jazelle state
- **Interrupt Disable bits**
  - I = 1: Disables IRQ
  - F = 1: Disables FIQ
- **T Bit**
  - T = 0: Processor in ARM state
  - T = 1: Processor in Thumb state
  - Introduced in Architecture 4T
- **Mode bits**
  - Specify the processor mode
- **New bits in V6**
  - **GE[3:0]** used by some SIMD instructions
  - **E** bit controls load/store endianness
  - **A** bit disables imprecise data aborts
  - **IT [abcde]** IF THEN conditional execution of Thumb2 instruction groups


THE ARCHITECTURE FOR THE DIGITAL WORLD®

8

ARM®

# Data alignment

- Prior to architecture v6 data accesses must be appropriately aligned for access size
  - Unaligned addresses will produce unexpected/undefined results



- Unaligned data can be accessed using **multiple** aligned accesses **combined** with shift/mask operations

# Exception Handling

- When an exception occurs, the core:
  - Copies CPSR into SPSR\_<mode>
  - Sets appropriate CPSR bits
    - Change to ARM state
    - Change to exception mode
    - Disable interrupts (if appropriate)
  - Stores the return address in LR\_<mode>
  - Sets PC to vector address
- To return, exception handler needs to:
  - Restore CPSR from SPSR\_<mode>
  - Restore PC from LR\_<mode>

0x1C	FIQ
0x18	IRQ
0x14	(Reserved)
0x10	Data Abort
0x0C	Prefetch Abort
0x08	Software Interrupt
0x04	Undefined Instruction
0x00	Reset

**Vector Table**

Vector table can also be at  
0xFFFF0000 on most cores

- Must be done in ARM state in most cores, but...  
...Thumb-2 capable cores can do this in Thumb state

# Introduction to Instruction Sets



## ARM Instruction Set

- All instructions are 32 bits long / many execute in a single cycle
- Instructions are conditionally executed
- A load / store architecture

- Example data processing instructions

```
SUB    r0,r1,#5
ADD    r2,r3,r3,LSL #2
ADDEQ  r5,r5,r6
```

```
r0 = r1 - 5
r2 = r3 + (r3 * 4)
IF EQ condition true r5 = r5 + r6
```

- Example branching instruction

```
B      <Label>
```

```
Branch forwards or backwards relative to
current PC (+/- 32MB range)
```

- Example memory access instructions

```
LDR    r0,[r1]
STRNEB r2,[r3,r4]
STMTFD sp!,{r4-r8,lr}
```

```
Load word at address r1 into r0
IF NE condition true, store bottom byte
of r2 to address r3+r4
Store registers r4 to r8 and lr on
stack. Then update stack pointer
```


# Thumb Instruction Set

- Thumb is a 16-bit instruction set
  - Optimized for code density from C code (~65% of ARM code size)
  - Improved performance from narrow memory
  - Subset of the functionality of the ARM instruction set
- Thumb is not a “regular” instruction set!
  - Constraints are not generally consistent
  - Targeted at compiler generation, not hand coding

# Thumb-2 Instruction Set

- Thumb-2 is a major extension to the Thumb ISA
  - Adds 32-bit instructions to implement almost all of the ARM ISA functionality
  - Retains the complete 16-bit Thumb instruction set
- Design objective: ARM performance with Thumb code density
  - No switching between ARM-Thumb states
  - Compiler automatically selects mix of 16 and 32 bit instructions

## Thumb 2 Performance / Density




## Processor Cores

**RealView®**  
Tools by ARM®

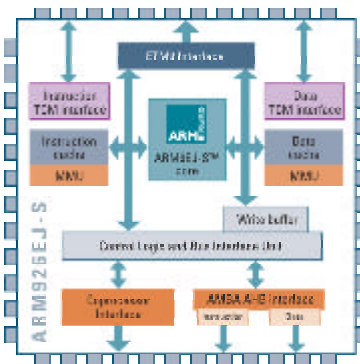


# ARM7TDMI Processor



- Architecture v4T
- 3-stage pipeline
- Single interface to memory

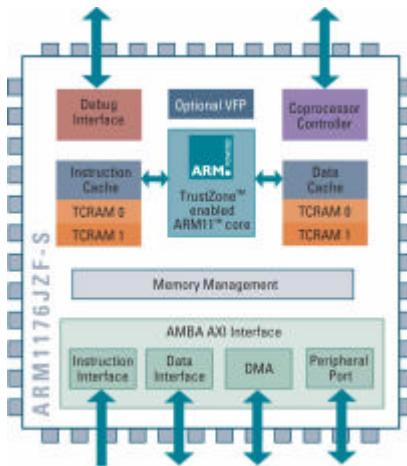
# ARM926EJ-S Processor



## ARM926EJ-S


- Architecture v5TE
- 5-stage pipeline
- Single-cycle 32x16 multiplier
- Caches and TCMs
- Memory management unit (MMU)
- 2 AHB memory interfaces
- Jazelle technology

# ARM1176JZ(F)-S Processor Core



- TrustZone
- 8-stage pipeline
- Branch prediction
- Four AXI memory ports
- IEM (Intelligent Energy Management)
- Integrated VFP coprocessor

# ARM11 MPCore Processor



- 1 – 4 MP11 processors
- Cache coherency
- Distributed interrupt controller

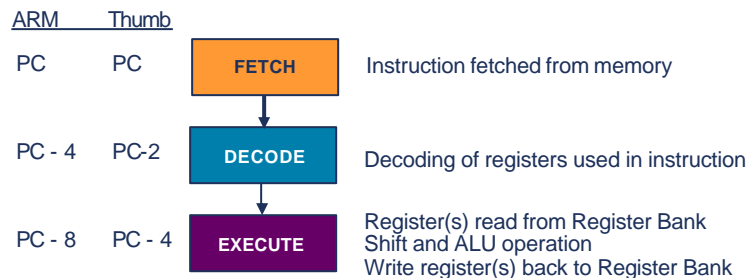


# The Instruction Pipeline



## The Instruction Pipeline

- The ARM7TDMI uses a 3-stage pipeline in order to increase the speed of the flow of instructions to the processor
  - Allows several operations to be performed simultaneously, rather than serially



- The PC points to the instruction **being fetched, not executed**
  - Debug tools will hide this from you
  - This is now part of the ARM Architecture and applies to all processors

# Optimal Pipelining

Cycle		1	2	3	4	5	6	7	8	9
Operation										
ADD		F	D	E						
SUB			F	D	E					
ORR				F	D	E				
AND					F	D	E			
ORR						F	D	E		
EOR							F	D	E	

F - Fetch D - Decode E - Execute

- All operations here are on registers (single cycle execution)
- In this example it takes 6 clock cycles to execute 6 instructions
- Clock cycles per Instruction (CPI) = 1

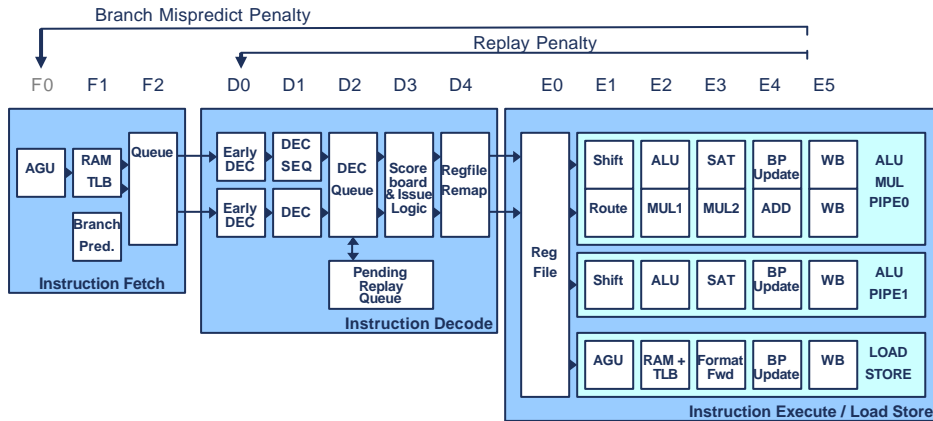
# Branch Pipeline Example

Cycle		1	2	3	4	5	6	7	8	9
Address Operation										
0x8000 BL 0x8FEC		F	D	E	E <sub>L</sub>	E <sub>A</sub>				
0x8004 SUB			F	D						
0x8008 ORR				F						
0x8FEC AND					F	D	E			
0x8FF0 ORR						F	D	E		
0x8FF4 EOR							F	D	E	

F - Fetch D - Decode E - Execute L - Linkret A - Adjust

- **Breaking the pipeline**
- Note that the core is executing in ARM state

# Cortex-A8 Integer Pipeline



- Optimising code to make use of the processor pipeline is very difficult
- Leave it to the compiler!!

## Reference Slides

# Reference Material

- **ARM ARM ("Architecture Reference Manual")**
  - ARM DDI 0100E covers v5TE DSP extensions
    - Can be purchased from booksellers - ISBN 0-201-737191 (Addison-Wesley)
    - Available for download from ARM's website
  - ARM v7-M ARM available for download from ARM's website
  - Contact ARM if you need a different version (v6, v7-AR, etc.)
- **Steve Furber "ARM system-on-chip architecture" - 2nd edition**
  - ISBN 0-201-67519-6 (Addison-Wesley)
- **Sloss, Symes & Wright – "ARM System Developer's Guide"**
  - ISBN: 1-55860-874-5 (Morgan Kaufman)
- **RVCT Assembler Guide**
  - Available for download from ARM's website
- **Technical Reference Manuals for processor core being used**
  - Available for download from ARM's website

# Naming Conventions

- **ARMx1z** (e.g. ARM710T) indicates cache & full MMU
- **ARMx2z** (e.g. ARM720T) indicates cache, MMU & Process ID support
- **ARMx3z** (e.g. ARM1136J-S) indicates physically mapped caches and MMU
- **ARMx4z** (e.g. ARM740T) indicates cache and MPU
- **ARMx5z** (e.g. ARM1156T2-S) indicates cache, MPU and error correcting memory
- **ARMx6z** (e.g. ARM966E-S) indicates write buffer but no caches
- **ARMx7z** (e.g. ARM1176JZ-S) indicates AXI bus, & physically mapped caches and MMU
- **ARMxy6** (e.g. ARM946E-S) indicates TCMs

# Which architecture is my processor?

Processor core	Architecture
<ul style="list-style-type: none"> <li>ARM7TDMI family <ul style="list-style-type: none"> <li>ARM720T, ARM740T</li> </ul> </li> </ul>	v4T
<ul style="list-style-type: none"> <li>ARM9TDMI family <ul style="list-style-type: none"> <li>ARM920T, ARM922T, ARM940T</li> </ul> </li> </ul>	v4T
<ul style="list-style-type: none"> <li>ARM9E family <ul style="list-style-type: none"> <li>ARM946E-S, ARM966E-S, ARM926EJ-S</li> </ul> </li> </ul>	v5TE, v5TEJ
<ul style="list-style-type: none"> <li>ARM10E family <ul style="list-style-type: none"> <li>ARM1020E, ARM1022E, ARM1026EJ-S</li> </ul> </li> </ul>	v5TE, v5TEJ
<ul style="list-style-type: none"> <li>ARM11 family <ul style="list-style-type: none"> <li>ARM1136J(F)-S</li> <li>ARM1156T2(F)-S</li> <li>ARM1176JZ(F)-S</li> <li>ARM11 MPCore</li> </ul> </li> </ul>	v6 v6 v6T2 v6Z v6
<ul style="list-style-type: none"> <li>Cortex family <ul style="list-style-type: none"> <li>ARM Cortex -A8</li> <li>ARM Cortex -R4(F)</li> <li>ARM Cortex -M3</li> <li>ARM Cortex -M1</li> </ul> </li> </ul>	v7-A v7-R v7-M v6-M
<ul style="list-style-type: none"> <li>For ARM processor naming conventions and features, please see the Appendix</li> </ul>	

## ARMv4T Cores:

	7TDMI	720T	740T	920T	940T	SA1100
Architecture	von Neumann	von Neumann	von Neumann	Harvard	Harvard	Harvard
Cache	None	8K Unified 4 words/line	8K Unified 4 words/line	16K Instr + 16K Data 8 words/line	4K Instr + 4K Data 4 words/line	16K Instr + 16K Data 4 words/line
Associativity	N/A	4-way	4-way	64-way	64-way	32-way
TCM	No	No	No	No	No	No
Replacement	N/A	Random	Random	Random Round Robin	Random	Round Robin
Write Strategy	N/A	Write Through	Write Through	Write Through Write Back	Write Through Write Back	Write Back
Write Buffer	None	8 Words 4 Addresses	8 Words 4 Addresses	16 Words 4 Addresses	8 Words 4 Addresses	8 Words 4 Addresses
MMU/MPU	None	MMU	MPU	MMU	MPU	MMU
Hi Vectors	No	Yes	No	Yes	Yes	Yes
Streaming	N/A	Yes	Yes	Yes	Yes	Yes
Standby Mode	No	No	No	Yes	Yes	Yes



## ARMv5 Cores:

	926EJ-S	946E-S	966E-S	968E-S	1026EJ-S	XScale
Architecture	Harvard	Harvard	Harvard	Harvard	Harvard	Harvard
Cache	4-128K Instr 4-128K Data 8 words/line	0-1024K Instr 0-1024K Data 8 words/line	None	None	0-128K Instr 0-128K Data 8 words/line	32K Instr 32K Data 8 words/line
Associativity	4-way	4-way	N/A	N/A	4-way	32-way
TCM	0-1024K Instr 0-1024K Data	0-1024K Instr 0-1024K Data	0-64M Instr 0-64M Data	0-64M Instr 0-64M Data	0-1024K Instr 0-1024K Data	No
Replacement	Random Round Robin	Random Round Robin	N/A	N/A	Random Round Robin	Random Round Robin
Write Strategy	Write Through Write Back	Write Through Write Back	N/A	Write Through Write Back	Write Through Write Back	Write Through Write Back
Write Buffer	16 Words 4 Addresses	16 Words Data or Address	12 Words Data or Address	12 Words Data or Address	8 Words Data or Address	8 x 16 Bytes Coalescing
MMU/MPU	MMU	MPU	None	None	MMU or MPU	MMU With extensions
Hi Vectors	Yes	Yes	Yes	Yes	Yes	Yes
Streaming	Yes	Yes	N/A	N/A	Yes	Yes
Standby Mode	Yes	Yes	Yes	Yes	Yes	Yes

## ARMv6 Cores:


	1136EJ(F)-S	1156T2(F)-S	1176JZ(F)-S	MPCore11
Architecture	Harvard	Harvard	Harvard	Harvard
Cache	4-64K Instr 4-64K Data 8 words/line	0-64K Instr 0-64K Data 8 words/line	4-64K Instr 4-64K Data 8 words/line	16-64K Instr 16-64K Data 8 words/line
Associativity	4-way	4-way	4-way	4-way
TCM	0-64K Instr 0-64K Data	0-256K Instr 0-256K Data	0-64K Instr 0-64K Data	None
Replacement	Random Round Robin	Random Round Robin	Random Round Robin	Random Round Robin
Write Strategy	Write Through Write Back	Write Through Write Back	Write Through Write Back	Write Through Write Back
MMU/MPU	MMU	MPU	MMU	MMU
Hi Vectors	Yes	Yes	Yes	Yes
Streaming	Yes	Yes	N/A	Yes
Standby Mode	Yes	Yes	Yes	Yes
Bus	AHB/APB	AXI	AXI	AXI
VFP Support	Yes	Yes	Yes	Yes

# Cortex Cores:

	Cortex-M3	Cortex-M1	Cortex-R4	Cortex-A8
Architecture	Harvard	Harvard	Harvard	Harvard
Cache	None	None	4-64K Instr 4-64K Data 8 words/line	16 or 32 Instr 16 or 32 Data 16 words/line
Associativity	N/A	N/A	4-way	4-way
TCM	None	0-1M Instr 0-1M Data	0-8M Instr 0-8M Data	None
Replacement	N/A	N/A	Random	Random
Write Strategy	N/A	N/A	Write Through Write Back	Write Through Write Back
MMU/MPU	MPU	None	MPU (optional)	MMU
Hi Vectors	No	No	Yes	Yes
Streaming	N/A	N/A	Yes	Yes
Standby Mode	Yes	Yes	Yes	Yes
Bus	AHB Lite/APB	AHB Lite/APB	AXI	AXI
VFP Support	No	No	Yes	Yes

# TrustZone Computing

- TrustZone adds a “parallel world” to allow trusted programs and data to be safely separated from the OS and applications
- Introduced for ARM1176, standard for ARMv7-A Cores
- Features:
  - New Secure Monitor Mode:** gate-keeper for secure state
  - New S-bit in CP15 to indicate when the processor is running in a secured state**
  - Security state exposed on external bus accesses to permit security-aware memory and peripherals**
  - Ability to restrict debug to non-secure state**



# NEON Media Processor Features

- Single Instruction Multiple Data (SIMD) Media Processor
- Targets audio and video codecs, image and speech processing, graphics, baseband processing, and general signal processing
- 3 Processing pipelines: Integer/fixed point, single precision floating point, IEEE vector floating point
- Efficient data handling
  - Best use of available memory bandwidth
  - Eliminates data arrangement overhead
  - Operates on separate register file
  - SIMD Framework excellent target for compilers



**End**